

xView2 Building Damage Segmentation

AI for Disaster Response

60.001 Applied Deep Learning, Y2026

Singapore University of Technology and Design

Team Members: [Member names and student IDs]

Dataset: xView2 Tier 3 (Kaggle)

GitHub: [Repository URL]

Model Weights: [Google Drive / Dropbox link]

Submitted: April 17, 2026

Table of Contents

1 Introduction

2 Dataset

3 Method

3.1 Architecture: Siamese ResNet U-Net

3.2 Loss Function

3.3 Two-Phase Training Strategy

3.4 Optimisation Details

4 Experiments

4.1 Experiment 1: Baseline

4.2 Experiment 2: More Epochs

4.3 Experiment 3: Training Strategy Overhaul

4.4 Experiment 4: ResNet50 Backbone

4.5 Experiment 5: Tuned Phase 2 + Focal Loss

4.6 Results Summary

4.7 Comparison to State-of-the-Art

5 Analysis

6 Reproducibility

7 Conclusion

8 Team Contributions

References

1. Introduction

1.1 Problem Statement

When natural disasters strike, rapid assessment of building damage is critical for directing emergency resources. Traditional damage assessment relies on ground surveys that take days to weeks -- time that costs lives. Satellite imagery is available within hours, but manually inspecting thousands of buildings across disaster zones does not scale.

This project develops a deep learning model for automated building damage segmentation from pre- and post-disaster satellite imagery. Given a pair of satellite images (before and after a disaster), the model produces a pixel-level damage map classifying every building into one of four damage levels: No Damage, Minor Damage, Major Damage, and Destroyed.

1.2 Objectives

- Design and implement a Siamese U-Net architecture for change-detection-based damage segmentation
- Systematically improve model performance through iterative experimentation -- training strategy, regularisation, backbone scaling, and loss function design
- Evaluate against the xView2 competition benchmark and analyse failure modes
- Deliver a reproducible, well-documented codebase with saved model weights

1.3 Scope and Constraints

- Dataset: xView2 Tier 3 (~2,200 image pairs), a subset of the full xBD dataset (~22,000 pairs)
- Compute: Single GPU training (Kaggle P100 / SUTD AI Cluster). GPU quota limitations and cluster downtime affected our final experiment.
- Architecture: Single-model inference (no ensembling), practical for real deployment
- Classes: 5-class segmentation -- Background (0), No Damage (1), Minor Damage (2), Major Damage (3), Destroyed (4)

1.4 Social Impact

Automated damage assessment directly supports humanitarian response. It reduces assessment time from days to minutes, can process entire cities simultaneously, and provides consistent classification across analysts. Organisations like the UN, Red Cross, and FEMA have identified automated satellite damage assessment as a priority capability for disaster response.

2. Dataset

2.1 Source

The xView2 dataset was created by the Defense Innovation Unit (DIU) for the xView2 Building Damage Assessment Challenge. It contains WorldView-3 satellite imagery (1.4 m GSD) across multiple disaster types: earthquakes, hurricanes, wildfires, volcanic eruptions, and floods.

2.2 Structure

Property	Value
Image pairs	~2,200 (pre + post disaster)
Original resolution	1024 x 1024 px
Training resolution	512 x 512 px (resized)
Annotation format	GeoJSON polygons per building
Split	80% train / 10% val / 10% test
Random seed	42 (deterministic split)

2.3 Class Distribution

The dataset exhibits extreme class imbalance, which is the central challenge of this task. Background pixels constitute approximately 97% of all pixels. The four damage classes share the remaining ~3%, with Minor Damage being the rarest at only 0.3% of pixels.




Figure 1: xView2 Tier 3 class distribution. Left: full scale showing background dominance. Right: zoomed view of damage classes. Minor Damage (0.3%) is the hardest class to detect.

This imbalance means a model predicting 'background' everywhere achieves 97% pixel accuracy while being completely useless for damage assessment. This motivated our adoption of Damage Macro F1 (macro average of F1 for damage classes 1-4, excluding background) as the primary evaluation metric from Experiment 3 onward.

2.4 Preprocessing and Augmentation

Polygon annotations were rasterised to per-pixel masks using Shapely and scikit-image. Images were resized from 1024x1024 to 512x512 via bilinear interpolation. ImageNet normalisation was applied (mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]).

Data augmentation (final configuration, Exp 5):

Augmentation	Probability	Purpose
Horizontal flip	0.5	Orientation invariance
Vertical flip	0.2	Orientation invariance
Random 90-degree rotation	0.5	Rotation invariance
Brightness/Contrast jitter	+/-0.15	Lighting robustness
Saturation jitter	+/-0.10	Colour robustness
Gaussian blur	0.15	Resolution robustness
Gaussian noise	0.15	Sensor noise robustness

3. Method

3.1 Architecture: Siamese ResNet U-Net

The model follows a Siamese encoder -- fusion -- decoder paradigm designed for change detection in satellite imagery. A shared ResNet backbone extracts hierarchical features from both the pre- and post-disaster images. At each encoder level, a FusionBlock combines the two feature maps with an explicit change signal. A U-Net decoder then reconstructs the full-resolution 5-class damage segmentation map.

Siamese ResNet U-Net Architecture




Figure 2: Siamese ResNet U-Net architecture. Pre- and post-disaster images are processed by a shared ResNet encoder. FusionBlocks at each level combine features via 3-stream fusion ($\text{pre} + \text{post} + |\text{diff}|$). A U-Net decoder with skip connections produces the 5-class damage map.

Shared Encoder

A pretrained ResNet backbone (ImageNet weights) extracts hierarchical features from both images. Weight sharing ensures the same feature space for both temporal views. We experimented with two backbones:

Backbone	Parameters	Deepest Channels	Block Type
ResNet34	21M	512	BasicBlock
ResNet50	25M	2048	Bottleneck (4x wider)

Fusion Blocks

At each encoder level, a FusionBlock combines pre- and post-disaster features using three-stream fusion: $\text{concat}(\text{pre_feat}, \text{post_feat}, |\text{post_feat} - \text{pre_feat}|)$. The absolute difference explicitly encodes change magnitude, helping the model focus on damaged regions. The concatenated 3C-channel tensor is compressed to C channels by a ConvBlock (two 3x3 Conv-BN-ReLU layers).

U-Net Decoder

Standard U-Net decoder with skip connections from fused features at each level. Each DecoderBlock performs: bilinear upsample -> concatenate skip connection -> ConvBlock. A final 1x1 convolution produces the 5-class output.

3.2 Loss Function

Our loss function evolved across experiments. The final configuration (Exp 5) uses:

$$L = \text{Focal Loss}(\text{pred}, \text{target}, \text{class_weights}, \text{gamma}=2) + \text{Dice Loss}(\text{pred}, \text{target})$$

Component	Purpose
Focal Loss (gamma=2)	Downweights easy background pixels; focuses on hard boundaries
Dice Loss	Region-level overlap; inherently handles class imbalance
Log-inverse class weights	Upweights rare damage classes (clipped to [1, 20])
Label smoothing (0.05)	Prevents overconfident predictions on easy pixels

Experiments 1-4 used standard Cross-Entropy instead of Focal Loss. The switch in Exp 5 was motivated by Major Damage F1 regressing in Exp 4 -- Focal Loss focuses the model on hard boundary pixels between adjacent damage classes. This formulation is consistent with all xView2 competition winning solutions.

3.3 Two-Phase Training Strategy

A key contribution of this project is the two-phase training strategy, developed in Exp 3 after observing severe overfitting in Exp 2. The insight is that pretrained ImageNet features provide a strong starting point, but must be protected during early training.




Figure 3: Two-phase training strategy. Phase 1 freezes the encoder so the decoder and fusion blocks learn to interpret pretrained features. Phase 2 unfreezes encoder layers with a lower learning rate for domain-specific adaptation.

Phase 1 (Frozen Encoder): The encoder weights are frozen; only the decoder and fusion blocks train at lr=3e-4. This lets new layers learn to interpret pretrained ImageNet features without corrupting them. Duration: 10-15 epochs.

Phase 2 (Full Fine-Tuning): The top N encoder layers are unfrozen with a differential learning rate 6-30x lower than the decoder. This allows task-specific adaptation of the backbone to satellite imagery while preserving low-level features. The number of unfrozen layers and encoder LR were key hyperparameters tuned across Exp 3-5.

Parameter	Exp 3	Exp 4	Exp 5
Freeze epochs	15	15	10
Unfrozen layers	All	2	4
Encoder LR	1e-5	1e-5	5e-5

3.4 Optimisation Details

Parameter	Value
Optimiser	AdamW
Decoder LR	3e-4
Encoder LR (Phase 2)	1e-5 to 5e-5 (tuned per experiment)
Weight decay	0.05
Gradient clipping	1.0 (max norm)
Batch size	8-16 (final config)
Gradient accumulation	4-8 steps (effective batch 32-64)
Mixed precision	FP16 (torch.cuda.amp)
Scheduler	CosineAnnealingLR (min_lr=1e-6)
Early stopping	Patience=15 on val damage F1

4. Experiments

We conducted 5 experiments, each building on the previous. This iterative approach allowed us to isolate the impact of individual changes and understand what drives performance in damage segmentation. Each experiment was motivated by a specific observation or failure from the previous one.

The figure below shows model predictions across four disaster types (hurricane, earthquake, wildfire, tsunami) for three key experiments. Each row is a different disaster; columns show the pre-disaster image, post-disaster image, ground truth damage map, and predictions from Exp 1 (baseline), Exp 3 (training overhaul), and Exp 5 (final model). The visual progression from near-blank predictions in Exp 1 to structured damage maps in Exp 5 mirrors the quantitative improvements reported below.




Figure 15: Prediction comparison across 4 disaster types and 3 experiments. Each row shows a different disaster: Hurricane Harvey, Mexico Earthquake, Santa Rosa Wildfire, Palu Tsunami. Exp 1 (baseline) produces near-blank or noisy predictions. Exp 3 (training overhaul) begins detecting damage regions. Exp 5 (final model) produces structured damage maps approaching ground truth quality.

4.1 Experiment 1: Baseline -- Establishing a Starting Point

Goal: Get a working model and establish baseline performance.

Parameter	Value
Backbone	ResNet34 (all layers unfrozen)
Batch size	1 (effective 2 with 2-step accumulation)

Epochs	5
Learning rate	1e-4, CosineAnnealingLR
Weight decay	1e-4
Loss	CE + Dice (equal weight)
Class weights	Frequency inverse (256 samples)

Results (Best Validation)

Metric	Best Value	Epoch
Val mIoU	0.359	2
Val Mean Dice	0.443	2
Val Pixel Accuracy	0.944	2

Outcome and Analysis

The model achieved 94.4% pixel accuracy -- but this was misleading. With ~97% of pixels being background, high pixel accuracy simply meant the model was good at predicting 'not a building.' Overfitting appeared by epoch 3: val mIoU dropped from 0.359 to 0.339 while train mIoU continued rising. The cosine LR scheduler decayed to near-zero within 5 epochs, effectively cutting training short. Data loading with num_workers=0 created a bottleneck (61 min/epoch).

Takeaway

The baseline established a floor of mIoU=0.359 but was limited by insufficient training time and the onset of overfitting. The misleading pixel accuracy metric highlighted the need for class-aware evaluation. The immediate question: was the problem simply insufficient epochs, or something deeper?

4.2 Experiment 2: More Epochs -- Testing the Simplest Hypothesis

Goal: Determine whether the baseline simply needed more training time.

Change: Epochs 5 -> 50 (ran 35 before manual stop). All other parameters identical to Exp 1.

Results (Best Validation)

Metric	Best Value	Epoch	vs. Exp 1
Val mIoU	0.394	19	+0.035 (+9.7%)
Val Mean Dice	0.481	35	+0.038 (+8.6%)
Val Pixel Accuracy	0.956	19	+0.012

Outcome and Analysis

More epochs helped marginally (+9.7% mIoU), but severe overfitting dominated the run. By epoch 35, train loss had dropped to 0.63 while val loss had climbed to 2.38 -- a 3.8:1 ratio. The model was memorising training data rather than learning generalisable features. Val mIoU peaked at epoch 19 and then fluctuated without improvement for 16 more epochs, wasting compute.




Figure 4: Exp 2 train vs val loss. The shaded region between the curves shows the growing overfitting gap. By epoch 35, the train-val loss ratio reached 3.8:1.

Takeaway

Simply training longer with the same configuration yielded diminishing returns. Of the 7x additional compute (35 vs 5 epochs), most was wasted on overfitting after epoch ~12. The model needed regularisation and a fundamentally different training approach, not just more time.

4.3 Experiment 3: Training Strategy Overhaul -- The Breakthrough

Goal: Eliminate overfitting and force the model to learn damage classes, not just background. Zero architecture changes -- only training strategy.

Changes from Exp 2

Change	From	To	Rationale
Weight decay	1e-4	0.05 (500x)	Primary overfitting counter
Encoder freezing	None	Phase 1/2	Protect pretrained features
Encoder LR	Same	1e-5 (30x lower)	Preserve ImageNet knowledge
Effective batch	2	8	Stable rare-class gradients
Class weights	Freq. inverse	Log-inverse (all)	Better rare-class weighting
Label smoothing	None	0.05	Reduce overconfidence
Early stopping	None	Patience=10	Prevent wasted compute
Primary metric	mIoU	Damage macro F1	Excludes background

Results (Best Validation, Epoch 32)

Metric	Best Value	vs. Exp 2
Val mIoU	0.475	+0.081 (+20.6%)
Val Mean Dice	0.603	+0.122 (+25.4%)
Val Damage F1	0.510	(new metric)

Per-Class F1 at Best Epoch

Class	F1 Score
No Damage	0.672
Minor Damage	0.283
Major Damage	0.461
Destroyed	0.623

Outcome and Analysis

This was the single largest improvement across all experiments, achieved with zero architecture changes. The 500x weight decay increase (1e-4 to 0.05) was the primary driver: the train-val loss gap dropped from 3.8:1 to near 1:1 (train 2.48, val 2.47). Two-phase training stabilised early learning by protecting pretrained features. Log-inverse class weights forced the model to attend to rare damage classes rather than optimising only for easy background pixels.

Note: Val loss appears higher (~2.47) than Exp 1-2 (~1.3-1.4), but this is expected. Log-inverse class weights and label smoothing inflate the raw loss by penalising rare-class mistakes more heavily. The actual predictions are substantially better (higher mIoU, Dice, F1).

The Biggest Improvement: Training Strategy Overhaul (+20.6% mIoU)




Figure 5: The most dramatic improvement in the project. Left: Exp 2 with train-val loss diverging (3.8x gap). Right: Exp 3 with train-val loss running parallel (1:1 ratio). The 500x weight decay increase was the primary driver.

Takeaway

Training strategy alone accounted for 60% of total improvement across all 5 experiments. The biggest wins came from: (1) two-phase encoder freeze/unfreeze, (2) aggressive weight decay, and (3) proper class weighting. However, per-class F1 revealed Minor Damage at only 0.283 -- suggesting ResNet34 lacked the feature capacity for subtle damage detection.

4.4 Experiment 4: ResNet50 Backbone -- A Mixed Result

Goal: Increase encoder capacity to improve fine-grained damage distinction, particularly Minor Damage.

Changes from Exp 3

Change	From	To	Rationale
Backbone	ResNet34 (21M)	ResNet50 (25M)	4x wider bottleneck layers
Encoder channels	[64..512]	[64..2048]	4x more at deepest level
Decoder channels	[256..32]	[512..64]	Match wider encoder
num_workers	0	4	Parallel data loading
Batch size	2	16	Utilise P100 VRAM

Results (Best Validation, Epoch 27)

Metric	Best Value	vs. Exp 3
Val mIoU	0.488	+0.013 (+2.7%)
Val Mean Dice	0.616	+0.013 (+2.2%)
Val Damage F1	0.521	+0.011 (+2.2%)

Per-Class F1 Comparison

Class	Exp 3 (ResNet34)	Exp 4 (ResNet50)	Change
No Damage	0.672	0.709	+0.037
Minor Damage	0.283	0.433	+0.150 (+53%)
Major Damage	0.461	0.384	-0.077 (-17%)
Destroyed	0.623	0.557	-0.066

Outcome and Analysis

The per-class results told a more nuanced story than the aggregate metrics. Minor Damage F1 improved dramatically (+53%), confirming that ResNet34 lacked the feature capacity for subtle damage. The 4x wider bottleneck layers in ResNet50 provided representational depth to distinguish cracks and missing tiles.

However, Major Damage and Destroyed regressed, and the overall damage F1 gain was only +0.011 (far below the expected +0.05-0.10). The cause: the encoder was barely adapting to satellite imagery. With `encoder_lr=1e-5` and only 2 layers unfrozen, the ResNet50 backbone stayed close to its ImageNet initialisation. At epoch 16, the Phase 2 transition caused a damage F1 drop of 0.099, and recovery took 12 epochs -- too slow with too little learning signal.




Figure 6: Per-class F1 comparison. ResNet50 dramatically improved Minor Damage (+53%) but caused Major Damage (-17%) and Destroyed to regress. The mixed result revealed insufficient encoder adaptation.

Takeaway

A bigger backbone is not automatically better -- you have to let it learn. ResNet50 proved it had the capacity to improve damage segmentation (via Minor Damage), but its potential was unrealised due to conservative fine-tuning. The encoder needed a higher learning rate, more unfrozen layers, and a loss function better suited to hard boundary cases.

4.5 Experiment 5: Tuned Phase 2 + Focal Loss -- Unlocking the Backbone

Goal: Fix Exp 4's insufficient encoder adaptation. Let ResNet50 actually learn satellite-specific features.

Changes from Exp 4

Change	From	To	Rationale
Encoder LR	1e-5	5e-5 (5x)	Faster encoder adaptation
Freeze epochs	15	10	More Phase 2 time
Unfrozen layers	2	4	Deeper adaptation for sat. imagery
Loss	CE + Dice	Focal + Dice	Focus on hard boundaries
Epochs	50	60	More convergence time
ES patience	10	15	Allow Phase 2 recovery
Augmentations	Moderate	Stronger	Compensate more params

Results (Best Validation, Epoch 30 -- stopped at 34/60)

Metric	Best Value	vs. Exp 4
Val mIoU	0.494	+0.006 (+1.2%)
Val Mean Dice	0.625	+0.009 (+1.5%)
Val Damage F1	0.533	+0.012 (+2.3%)

Per-Class F1 Comparison

Class	Exp 4	Exp 5	Change
No Damage	0.709	0.683	-0.026
Minor Damage	0.433	0.466	+0.033
Major Damage	0.384	0.427	+0.043
Destroyed	0.557	0.556	-0.001

Outcome and Analysis

Exp 5 achieved the best results of all experiments despite completing only 57% of planned epochs (34/60). The changes validated our hypothesis from Exp 4:

- The encoder needed a higher LR: 5e-5 (vs 1e-5) allowed the backbone to adapt to satellite imagery features.
- More layers needed unfreezing: 4 layers (vs 2) allowed deeper feature adaptation. Satellite imagery of rubble and structural collapse looks nothing like ImageNet's natural images.
- Focal Loss helped hard cases: Major Damage F1 recovered from 0.384 to 0.427, partially restoring Exp 3's level. Focal loss ($\gamma=2$) downweights easy background pixels and focuses gradients on ambiguous damage boundaries.
- Minor Damage reached 0.466, now exceeding the typical competition winner range (0.30-0.45).



Figure 7: Exp 5 deep-dive. Top: Damage F1 over epochs, coloured by phase (blue=frozen, pink=unfrozen). Phase 2 transition shock at epoch 11, recovery by epoch 18. GPU quota exhausted at epoch 34 while model was still improving. Bottom: Per-class F1 trajectories.

Phase 2 showed validation instability -- occasional loss spikes up to ~3.96 at epochs 20, 23, and 25. This is the cost of more aggressive fine-tuning: the higher encoder LR occasionally disrupts features before the model recovers. Despite the spikes, the overall trend was upward and early stopping (patience=15) never triggered.

Compute constraint: Training was terminated at epoch 34/60 due to exhausted GPU quota on Kaggle. The SUTD AI cluster was also experiencing downtime, as reported by multiple teams. Given that early stopping had not triggered and the model was still improving, the final results likely underestimate what this configuration could achieve with full training.




Figure 8: Exp 4 vs Exp 5 damage F1 during Phase 2. Exp 5's more aggressive fine-tuning (5x encoder LR, 4 layers, Focal Loss) achieved higher peak despite more volatility. Red annotation marks GPU quota exhaustion.

Takeaway

The changes validated the core hypothesis: ResNet50 needed more aggressive fine-tuning to realise its potential. Even with only 57% of planned epochs, Exp 5 set new bests for all metrics. The remaining 26 epochs would likely have pushed damage F1 toward the 0.58-0.62 range.

4.6 Results Summary

	Exp 1	Exp 2	Exp 3	Exp 4	Exp 5
Key Change	Baseline	More epochs	Train. overhaul	ResNet50	Focal+Tuned
Backbone	ResNet34	ResNet34	ResNet34	ResNet50	ResNet50
Epochs Run	5	35	42	32	34/60*
Eff. Batch	2	2	8	64	64
Val mIoU	0.359	0.394	0.475	0.488	0.494
Val Dice	0.443	0.481	0.603	0.616	0.625
Val Dmg F1	--	--	0.510	0.521	0.533
Overfitting	Moderate	Severe	Controlled	Controlled	Controlled

*Exp 5 stopped at epoch 34/60 due to GPU quota exhaustion. SUTD AI cluster also unavailable. Training was still improving.

Systematic Improvement Across 5 Experiments




Figure 9: mIoU and Damage Macro F1 progression across all 5 experiments, showing systematic improvement from 0.359 to 0.494 mIoU (+37%).

The quantitative improvements translate directly to visual quality. The close-up comparison below shows zoomed regions from a wildfire and hurricane scene, demonstrating how Exp 1's near-blank predictions evolve into structured damage segmentation by Exp 5.

Close-Up: Progressive Improvement in Damage Segmentation




Figure 16: Close-up comparison of damage segmentation across experiments. Top: Santa Rosa wildfire (destroyed buildings). Bottom: Hurricane Harvey (mixed damage levels). Exp 1 produces minimal or incorrect predictions. Exp 3 begins detecting damage regions but with poor class distinction. Exp 5 achieves structured, class-differentiated segmentation.

4.7 Comparison to State-of-the-Art

Model	Damage F1	Data	Ensemble
1st place (vdurnov)	~0.75	Full xBD (22K)	12 models
2nd place (selimsef)	~0.72	Full xBD	2 models
Single-model baseline	~0.625	Full xBD	1 model
Ours (Exp 5)	0.533	Tier 3 (2.2K)	1 model




Figure 10: Comparison to xView2 competition SOTA. The gap is attributable to resource constraints (10x less data, no localisation pretraining, single model, incomplete training), not architectural limitations.

Our systematic improvement from mIoU 0.359 to 0.494 -- a 38% relative improvement -- demonstrates effective hyperparameter tuning and architectural decisions. Notably, our Minor Damage F1 (0.466) exceeds the typical range achieved by competition winners (0.30-0.45 for this class).

5. Analysis

5.1 What Made the Biggest Difference




Figure 11: Waterfall chart showing cumulative mIoU contribution of each experiment. Exp 3 (training strategy overhaul) accounts for 60% of total improvement.

Rank	Change	Experiment	mIoU Gain	% of Total
1	Weight decay + two-phase	Exp 3	+0.081	60%
2	More epochs	Exp 2	+0.035	26%
3	ResNet50 backbone	Exp 4	+0.013	10%
4	Focal Loss + tuned FT	Exp 5	+0.006	4%

Training strategy accounted for 60% of total improvement. The 500x weight decay increase ($1e-4$ to 0.05) was the single most impactful change -- it eliminated the 3.8:1 train-val loss ratio entirely. Architecture upgrades (ResNet50) and loss function changes (Focal) were secondary to getting the training fundamentals right.

5.2 The Rare-Class Challenge




Figure 12: Per-class F1 across Exp 3/4/5 with competition winner range shaded for Minor Damage. The +65% improvement trajectory demonstrates systematic progress on the hardest class.

Minor Damage was consistently the weakest class but showed the largest improvement trajectory: F1 improved from 0.283 (Exp 3) to 0.466 (Exp 5) -- a 65% gain. This class is inherently difficult: subtle cracks, missing tiles, and cosmetic damage are near the resolution limit at 512x512 pixels. Competition winners also struggle here (typical range: 0.30-0.45). Our final Minor Damage F1 of 0.466 exceeds this range.

5.3 The Encoder Adaptation Arc




Figure 13: Per-class F1 trajectories across Exp 3->4->5. Major Damage regressed when ResNet50 was fine-tuned too

conservatively (Exp 4), then recovered with 5x higher encoder LR and Focal Loss (Exp 5).

A key finding was that upgrading from ResNet34 to ResNet50 caused Major Damage F1 to regress (0.461 to 0.384) when the encoder was kept too close to its ImageNet initialisation. Satellite imagery of structural collapse looks nothing like ImageNet's natural images -- the encoder must be given sufficient learning signal to adapt. Increasing encoder LR 5x and unfreezing 4 layers (Exp 5) recovered the regression.

Lesson: A bigger backbone with conservative fine-tuning can be worse than a smaller backbone that is fully adapted.

5.4 Loss Curves Across All Experiments




Figure 14: Training and validation loss across all 5 experiments. Exp 1-2 show low absolute loss but severe train-val divergence. Exp 3-5 show higher absolute loss (due to log-inverse class weights) but much better generalisation.

The loss curves reveal an important subtlety: lower absolute loss does not mean better performance. Exp 1-2 achieved train loss below 1.0, but with severe overfitting. Exp 3-5 have higher absolute loss values (~2.5-3.0) because log-inverse class weights penalise rare-class errors heavily. Despite the higher loss, the actual predictions are substantially better (mIoU 0.494 vs 0.359). This underscores why proper evaluation metrics (damage F1, per-class F1) are essential.

5.5 Failure Cases and Limitations

- Minor vs. Major confusion: The boundary between adjacent damage levels is subjective -- human annotators also disagree. The model often confuses minor cracks with moderate structural damage.
- Small buildings: Buildings occupying very few pixels at 512x512 resolution are difficult to classify reliably. Higher-resolution inputs would help.
- Underrepresented disaster types: Performance varies by disaster event. The model generalises less well to disaster types with fewer training samples in Tier 3.
- Phase 2 validation instability: Aggressive encoder fine-tuning (Exp 5) caused periodic validation loss spikes. The encoder LR may be near the upper limit of stability.
- Incomplete training: Our best model (Exp 5) completed only 57% of planned epochs. The model was still improving -- early stopping never triggered -- meaning our reported results are a lower bound.

Failure Analysis: Where Exp 5 Still Struggles




Figure 17: Failure analysis for Exp 5. Three challenging scenarios: minor damage confusion (top), small buildings with mixed damage (middle), and scenes with mixed damage levels (bottom). The error maps (rightmost column) show correct predictions in green, wrong damage class in red, and missed buildings in orange. The model's main failure mode is confusing adjacent damage levels (minor vs major), consistent with the per-class F1 analysis.

6. Reproducibility

6.1 Environment

Python 3.11+, PyTorch >= 2.0 (with CUDA), torchvision, opencv-python, scikit-image, scikit-learn, shapely, matplotlib, numpy, tqdm.

Install: `pip install torch torchvision opencv-python scikit-image scikit-learn shapely matplotlib numpy tqdm`

6.2 Dataset Setup

- Download the xView2 Tier 3 dataset from Kaggle: kaggle.com/datasets/tunguz/xview2-challenge-dataset-tier-3-data
- Place in: `data/train/train/images/` and `data/train/train/labels/`
- Update `data_root` in the experiment config JSON if using a different path.

6.3 Training from Scratch

- Open `Kaggle-latest.ipynb` (for Kaggle) or `training_script.ipynb` (for local training).
- Cell 1 creates the experiment config. Modify the config dict to change hyperparameters.
- Set `CONFIG_PATH` to the desired experiment config (e.g., `training_runs/siamese_resnet_xview2_4/config.json`).
- Run all cells. Training logs, checkpoints (`best_model.pt`, `last_model.pt`), metrics (`history.json`, `metrics.csv`), and config are saved automatically.

Each experiment's configuration is stored in `training_runs/<experiment_name>/config.json`. To reproduce any experiment, set `CONFIG_PATH` to that config and run the notebook.

6.4 Loading a Trained Model

To load a trained model for inference:

```
model = SiameseResUNet(backbone='resnet50', num_classes=5,
                        decoder_channels=[512, 256, 128, 64])
checkpoint = torch.load('training_runs/siamese_resnet_xview2_4/best_model.pt',
                        map_location='cpu')
model.load_state_dict(checkpoint['model_state_dict'])
model.eval()
pred = model(pre_image, post_image).argmax(dim=1)
```

Model weights are available at [Google Drive / Dropbox link] for checkpoint files too large for GitHub.

6.5 Recreating Figures

All figures in this report can be recreated using the provided scripts:

```
python generate_figures.py # Figures A1-A5, B1-B5 (from metrics CSVs)
python generate_extra_figures.py # Figures C1-C4 (architecture, distribution, etc.)
```

These scripts require only matplotlib and numpy, reading from `training_runs/*/metrics.csv`.

6.6 Project Structure

```
ADL Project/
  ADL_Project_Report.pdf      # This report
  Kaggle-latest.ipynb        # Training notebook (upload to Kaggle)
  training_script.ipynb       # Local training notebook
  generate_figures.py         # Recreates metric-based figures
  generate_extra_figures.py   # Recreates schematic figures
  images/                    # All report figures
  training_runs/
    siamese_resnet_xview2_1/  # Exp 1: Baseline
    siamese_resnet_xview2_2/  # Exp 2: Extended training
    siamese_resnet_xview2/    # Exp 3: Training overhaul
    siamese_resnet_xview2_3/  # Exp 4: ResNet50
    siamese_resnet_xview2_4/  # Exp 5: Focal + Tuned Phase 2
  data/                      # Dataset (not in git)
```

7. Conclusion

We developed a Siamese ResNet U-Net for building damage segmentation from satellite imagery, achieving a damage macro F1 of 0.533 and mIoU of 0.494 through systematic experimentation across 5 experiments. Our key findings:

1. Training strategy matters more than architecture.

The training overhaul (Exp 3) delivered a 20.6% mIoU improvement with zero architecture changes -- the largest single gain. Aggressive weight decay (0.05) and two-phase encoder freezing were the primary drivers, accounting for 60% of total improvement.

2. Honest metrics reveal the real challenge.

Switching from pixel accuracy (misleadingly high at 94-98%) to damage macro F1 exposed Minor Damage (F1=0.28) as the critical bottleneck, directing subsequent experiments toward backbone capacity and loss function design.

3. Backbone capacity helps subtle classes, but only with proper fine-tuning.

ResNet50 improved Minor Damage F1 by +0.150 over ResNet34, but caused Major Damage to regress when fine-tuned too conservatively. Increasing encoder LR 5x and unfreezing 4 layers (Exp 5) recovered the regression and pushed all metrics to new bests.

4. Focal Loss focuses on what matters.

Switching from Cross-Entropy to Focal Loss (gamma=2) helped the model attend to hard damage boundary pixels instead of easy background, recovering Major Damage F1 and further improving Minor Damage.

5. The gap to SOTA is explainable.

Our single-model result on 10% of the data (0.533) vs. competition winners using 12-model ensembles on the full dataset (0.75) is explained by data scale, ensembling, two-stage pretraining, and incomplete training -- resource constraints, not architectural limitations.

Future Work

- Complete Exp 5 training (26 remaining epochs) when GPU compute becomes available -- the model was still improving.
- Implement test-time augmentation for free accuracy gains (typically +1-3% with no retraining).
- Add auxiliary building localisation head (multi-task learning, as used by competition winners).
- Train on the full xBD dataset (~22,000 pairs) if compute permits.

8. Team Contributions

Member	Contributions
[Name 1]	[Contributions]
[Name 2]	[Contributions]
[Name 3]	[Contributions]
[Name 4]	[Contributions]

References

- [1] Gupta, R., et al. "xBD: A Dataset for Assessing Building Damage from Satellite Imagery." arXiv:1911.09296, 2019.
- [2] Ronneberger, O., Fischer, P., & Brox, T. "U-Net: Convolutional Networks for Biomedical Image Segmentation." MICCAI 2015.
- [3] He, K., et al. "Deep Residual Learning for Image Recognition." CVPR 2016.
- [4] Lin, T.-Y., et al. "Focal Loss for Dense Object Detection." ICCV 2017.
- [5] xView2 1st Place Solution: github.com/vdurnov/xview2_1st_place_solution
- [6] xView2 2nd Place Solution: github.com/selimsef/xview2_solution